


Generative Adversarial Network

Instructor: Seunghoon Hong

Course logistics

- New assignment will be out today
 - **Deadline:** 23:59:59 November 22th
 - **Quiz:** Nov. 27th

Recap: objective of generative models



$$\theta^* = \arg \max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} \log p_{\theta}(x)$$

Recap: Autoregressive models

- Explicit optimization of likelihood based on chain rule

$$\log p_{\theta}(x) = \sum_{t=1}^d \log p_{\theta}(x^t | x^1, \dots, x^{t-1})$$

For d-dimensional data $x = (x^1, x^2, \dots, x^d)$

- Advantages:
 - Optimizing exact likelihood
- Disadvantages:
 - Sequential generation process ($O(n)$)
 - No latent variable to control the generation process

Recap: Variational Autoencoder

- Explicit optimization of variational lower-bound of likelihood

$$\log p(x) \geq \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - D_{KL}(q_{\phi}(z|x)||p(z))$$

- Advantage:
 - Latent variable model (i.e. we can control generation via z)
 - Feedforward generation
- Disadvantages:
 - Lower-bound optimization (there is a gap between actual likelihood and lower-bound if $q \neq p$).
 - Generally not satisfactory generation quality

Explicit density models

- Explicitly modeling the likelihood of data
- Autoregressive models and VAEs are both explicit density models

Autoregressive models

$$\log p_{\theta}(x) = \sum_{t=1}^d \log p_{\theta}(x^t | x^1, \dots, x^{t-1})$$

Variational Autoencoder

$$\log p(x) \geq \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - D_{KL}(q_{\phi}(z|x) || p(z))$$

Challenges in explicit density models

- What if we cannot measure the likelihood?
- Example: conditional generation (e.g. machine translation)
 - In some cases we do not have a paired data
 - Example: unaligned data for machine translation → *pair가 없어 likelihood measure 불가능.*
 - In this case, we cannot measure the conditional probability

Implicit density model

- Modeling distribution without explicit likelihood estimation

Today's agenda

- Generative Adversarial Network (GAN)

Generation task as an adversarial game


(위조지폐)

- Intuitive example: a game of counterfeiting

Adversarial game
between two players



VS



Goal of counterfeiter

Make fake money as realistic as possible

경찰 속일 정도로 만들면 충분

Generator

즉, 이 위조지폐의
quality보다는
경찰 속이기에 집중.

위조지폐 탐지할 정도의 집중

Goal of police officer

Detect counterfeit money from real one


Discriminator

→ End: Equilibrium (50%)

Note: money의 quality는
측정하지 않음 []

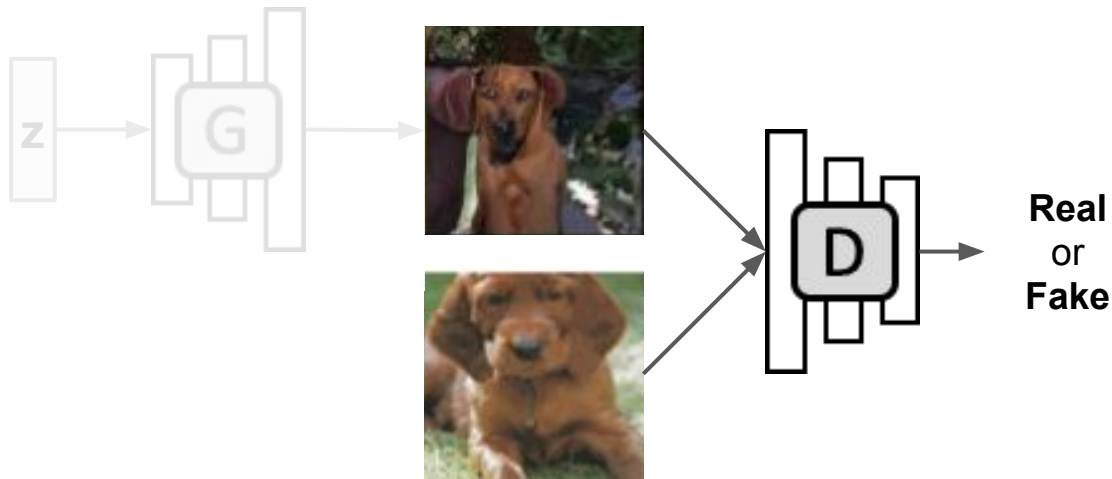
Generative Adversarial Network (GAN)

- Learning to generate via minimax optimization



Generative Adversarial Network (GAN)

- Learning to generate via minimax optimization




- **Discriminator (D)**
tries to tell if
its input is real or fake

$z \sim p_z$, binary
classification.

Generative Adversarial Network (GAN)

- Learning to generate via minimax optimization



- **Discriminator (D)**
tries to tell if
its input is real or fake
- **Generator (G)**
tries to fool
the discriminator

Learning objective

- Minimax objective function

$$\min_{\theta_g} \max_{\theta_d} \left[\underbrace{\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x)}_{\text{Discriminator output for real input } x} + \underbrace{\mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))}_{\text{Discriminator output for the generator output } G(z)} \right]$$

Given a generator $G(z; \theta_g)$

- Discriminator tries to maximize the objective such that $D(x)$ is close to 1 (real) and $D(G(z))$ is close to 0 (fake) \rightarrow solves binary classification problem

Given a discriminator $D(x; \theta_d)$

- Generator tries to minimize objective such that $D(G(z))$ is close to 1 (discriminator is fooled into thinking generated $G(z)$ is real)

Learning objective

- Minimax objective function

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. **Gradient descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

Optimization challenge

Alternate between:

1. Gradient ascent on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. Gradient descent on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$


Gradient on the generator

is dominated by the region

where the sample is already good

When sample is likely fake, want to learn from it to improve generator. But gradient in this region is relatively flat!

Gradient signal dominated by region where sample is already good



Optimization challenge

Alternate between:


1. Gradient ascent on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. Gradient ascent on generator but different using different objective

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$

Also known as **non-saturating loss**



GAN training algorithm

for number of training iterations **do**

for k steps **do**

- \approx batch size $2m$.
- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
 - Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
 - Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D_{\theta_d}(x^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(z^{(i)}))) \right]$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by ascending its stochastic gradient (improved objective):

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(z^{(i)})))$$

end for

what it means?

The very first GAN results

Get sharp image.

Goodfellow et al., Generative adversarial networks, In NIPS, 2014





Synthesized images by the generator

Nearest neighbors in a training set
(of the rightmost generated example)

The very first GAN results

Goodfellow et al., Generative adversarial networks, In NIPS, 2014



Case study: DCGAN


- Techniques to improve GAN training

Architecture guidelines for stable Deep Convolutional GANs


- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

Case study: DCGAN

- Generator architecture



Generated images - LSUN bedroom dataset



Radford et al., unsupervised representation learning with deep convolutional generative adversarial networks, In ICLR, 2016

Sample interpolation

interpolate.

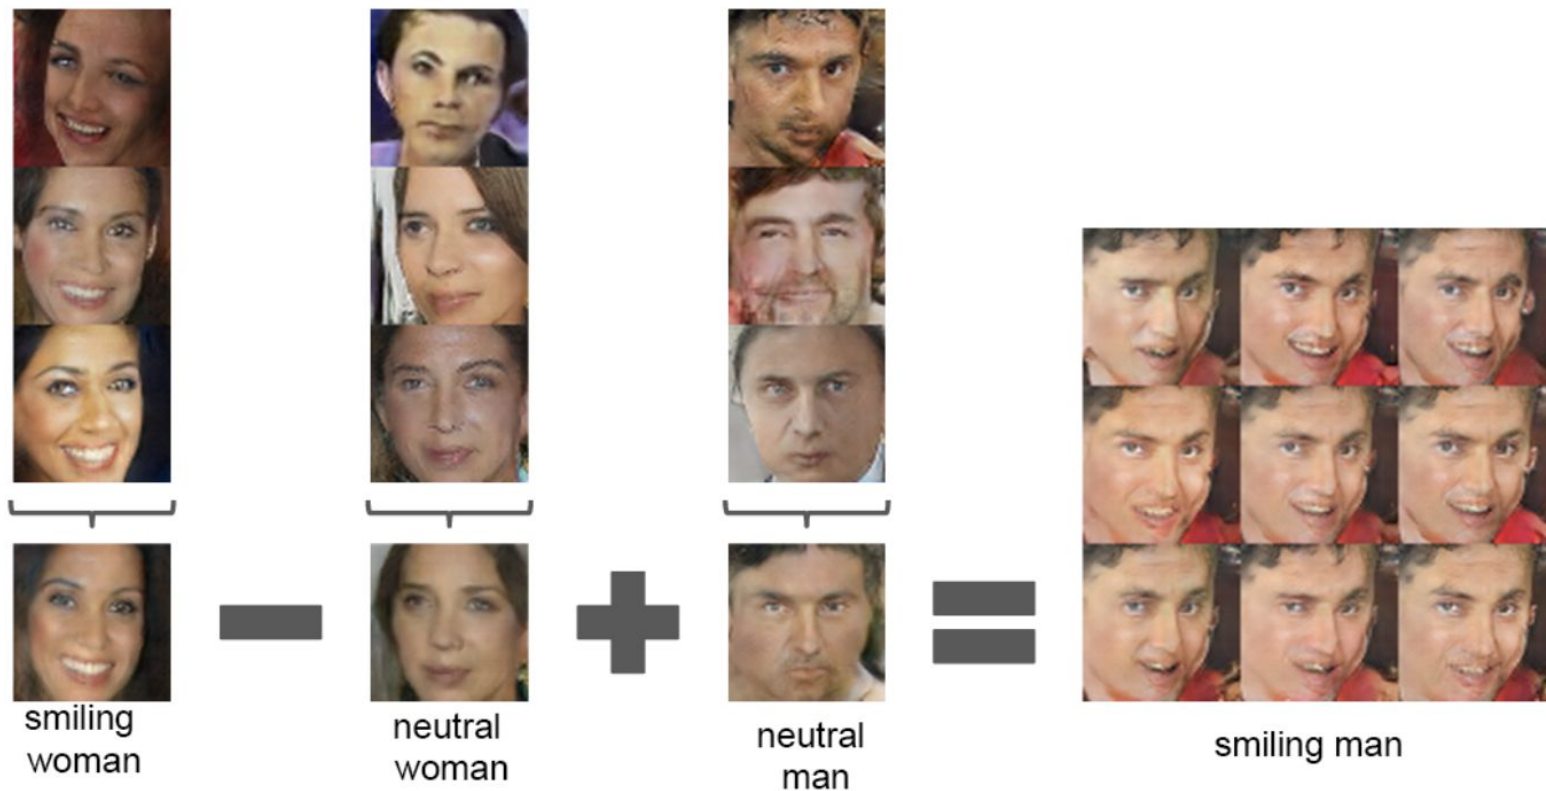
z_2
↓



Sample interpolation




Arithmetic on latent variable



Radford et al., unsupervised representation learning with deep convolutional generative adversarial networks, In ICLR, 2016

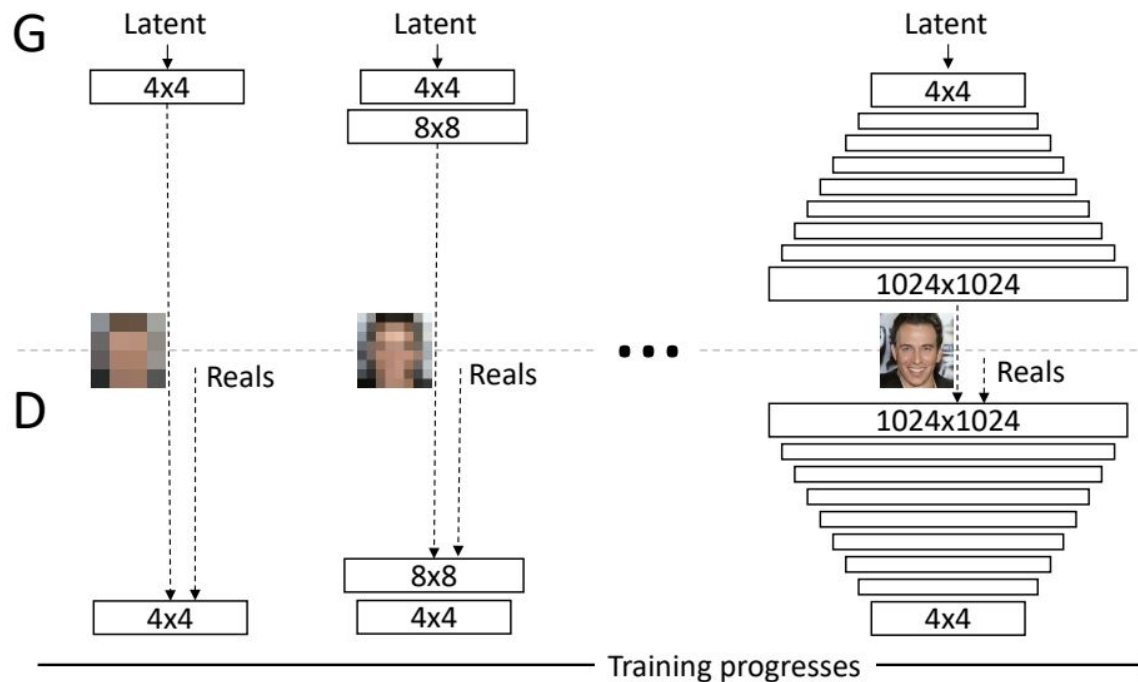
Arithmetic on latent variable



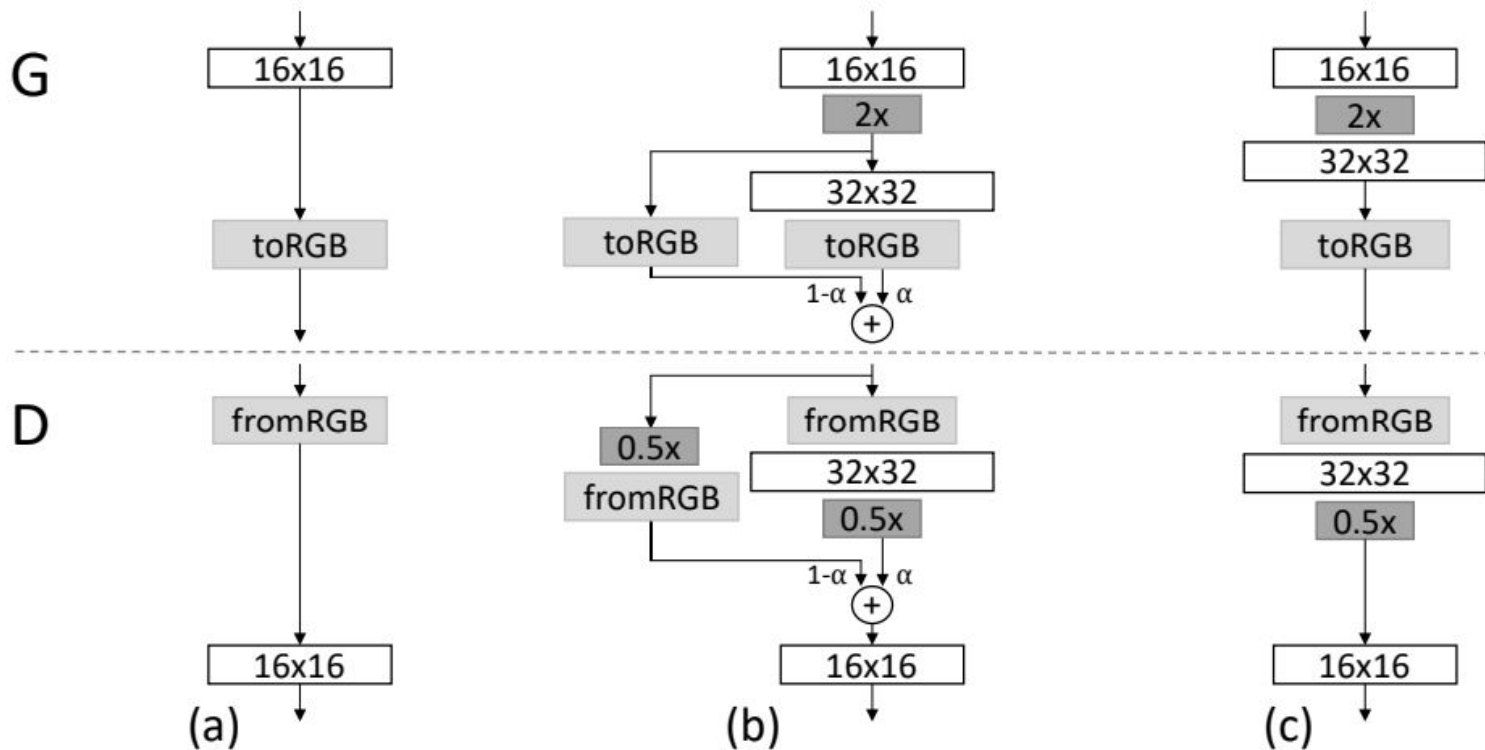
Radford et al., unsupervised representation learning with deep convolutional generative adversarial networks, In ICLR, 2016

Progressive growing of GAN

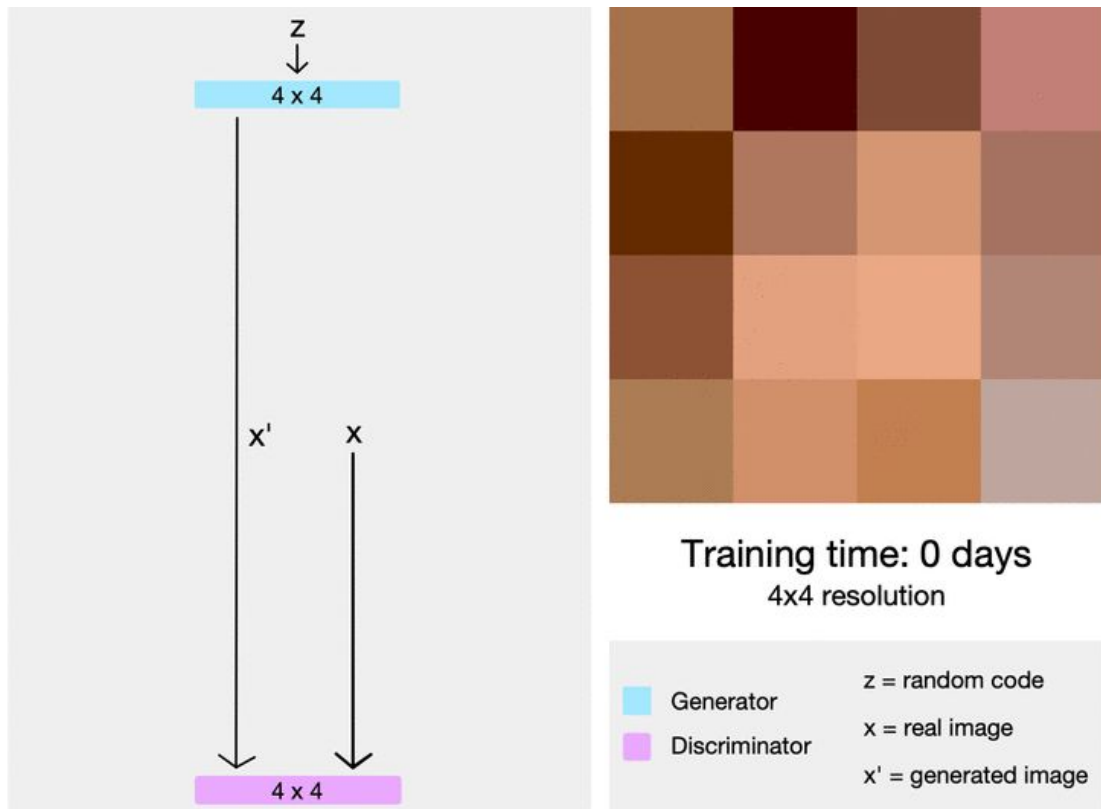
- Improve the generation quality through hierarchical generation




Progressive growing of GAN



Progressive growing of GAN

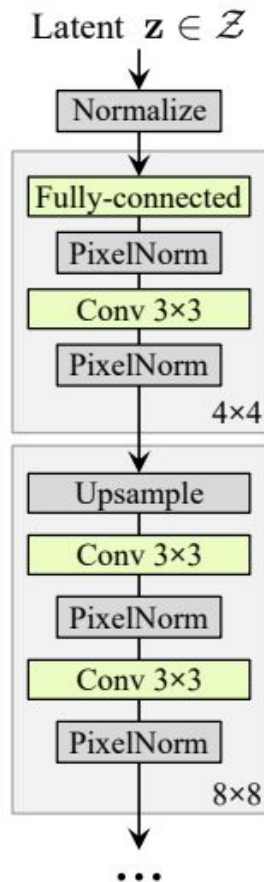


BigGAN

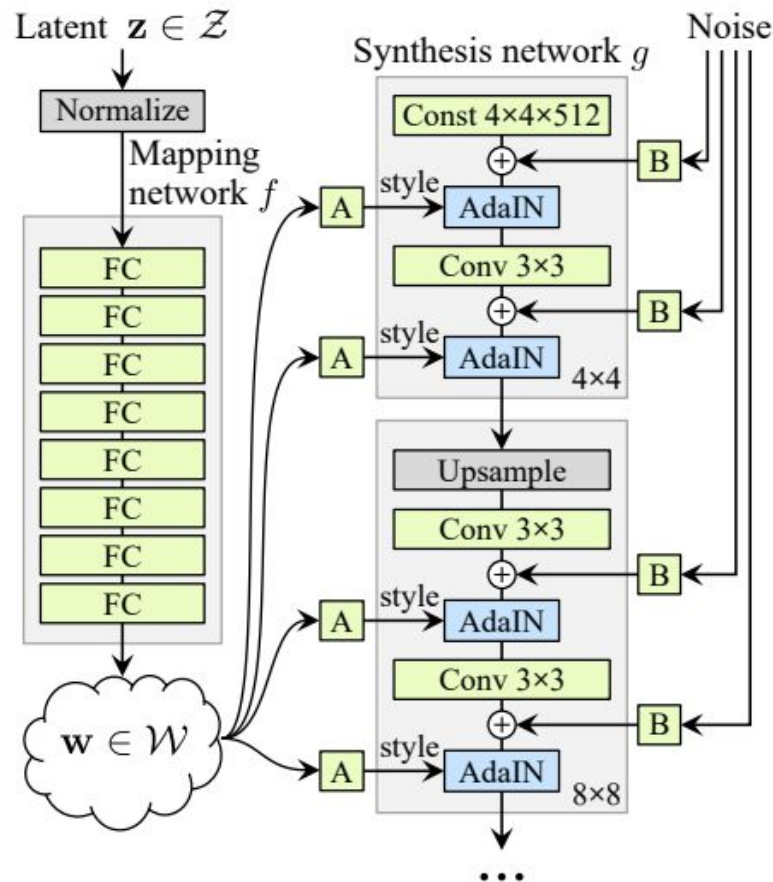


StyleGAN

- Improved version of progressive GAN
- Deep embedding layers of latent variable
- Injecting latent variable via modulation



(a) Traditional



(b) Style-based generator

StyleGAN: results



StyleGAN: sample interpolation



Keras et al., A Style-Based Generator Architecture for Generative Adversarial Networks, In CVPR, 2019

StyleGAN: try yourself

<https://thispersondoesnotexist.com/>



Challenges in GAN

- Is adversarial game stable?



Challenges in GAN

- Is adversarial game stable? → it is turned out to be not
- One representative problem: **mode-collapse**

(It does not happen)


combine modes. blurred image

$$0.9 \frac{0}{1}$$

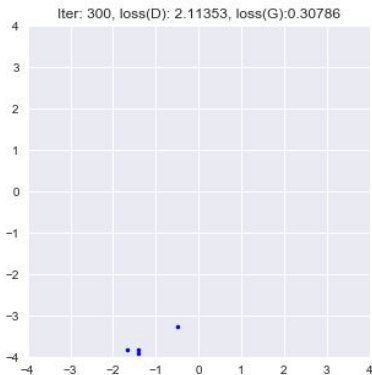
in

autoencoder

True Data



GAN output



When mode-collapse happens,
the generator models **only part of** the true
data distribution (e.g. one data mode)

Why does it happens?

→ intuitively, a single mode is also
indistinguishable from the real data