

Attention and versatile architectures

Presenter: Seunghoon Hong

Course overview

- ❑ Image classification
- ❑ Object detection
- ❑ Semantic segmentation
- ❑ Visualization
- ❑ Style transfer
- ❑ Adversarial attacks
- ❑ Text modeling
- ❑ Machine translation
- ❑ Image captioning
- ❑ Visual question answering
- ❑ Image generation
- ❑ Text generation
- ❑ Img-to-img translation
- ❑ Attention and versatile networks
- ❑ Self- and Semi-supervised learning
- ❑ Multi-modal learning
- ❑ Graph neural networks



So far, we learned various neural networks

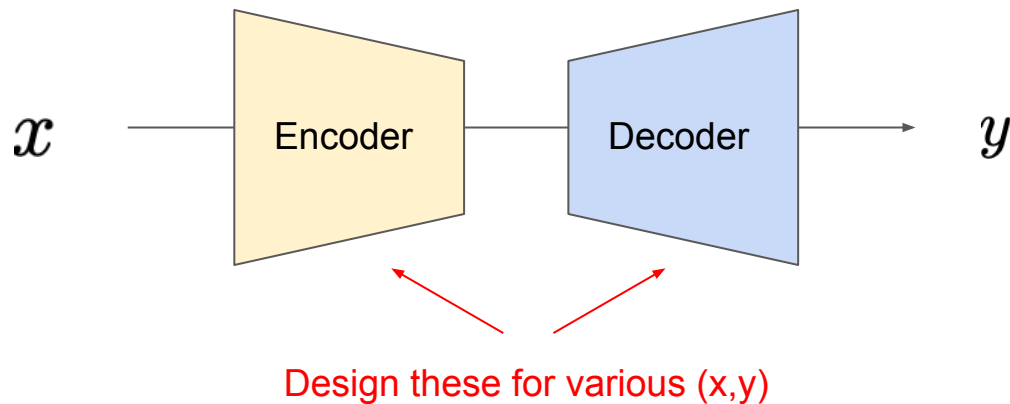
- Multi-Layer Perceptron (MLP)
 - Linear layer + Non-linear function
- Convolutional Neural Network (CNN)
 - Convolution + Pooling + Non-linear function + MLP
- Recurrent Neural Network (RNN)
 - MLPs for recurrent update + gating functions
- Attention and Transformer
 - Dot-product attention + MLPs

We also learned about various tasks

- Classification
- Detection
- Segmentation
- Text encoding
- Machine translation
- Image captioning
- Visual question answering
- ...


We are ready to be DL practitioners

- Let's try solving some problems!




Practice 1: speaker identification

Given a speech signal, identify the speaker given the list of known people.




X: Mel-spectrogram




Y: {Yumi, Jack, **John**, Jane}

Practice 2: speech to text translation

Given a speech signal, transcribe the content into the text



X: Mel-spectrogram




“Siri, turn off the light”

Y: natural language

Practice 3: creating music from score

Given the music score, create the piano song playing the score




X: music score

Y: Mel-spectrogram

Problem 4: creating music and lyrics from score

Given the music score, create the piano song playing the score and lyrics



Different data (tasks), different models

- Why do we have different models for different data and tasks?
- What happens if we apply different combinations of data and models?
 - Example 1: CNN for sequences.
 - Example 2: RNN for images.
 - Example 3: MLP for images.

Inductive bias in neural networks

- **Inductive bias (formal definition):** a set of assumptions that the learner uses to predict outputs given input that has not observed

Inductive bias in neural networks

- **Inductive bias (formal definition):** a set of assumptions that the learner uses to predict outputs given input that has not observed
- Many neural network architectures rely on certain type of inductive biases
 - Mostly depends on locality and invariance of data
 - Example1: Convolution exploits 2D grid data structure and translation invariance
 - Example 2: RNN exploits 1D chain data structure and first-order Markov assumption

Inductive bias in neural networks

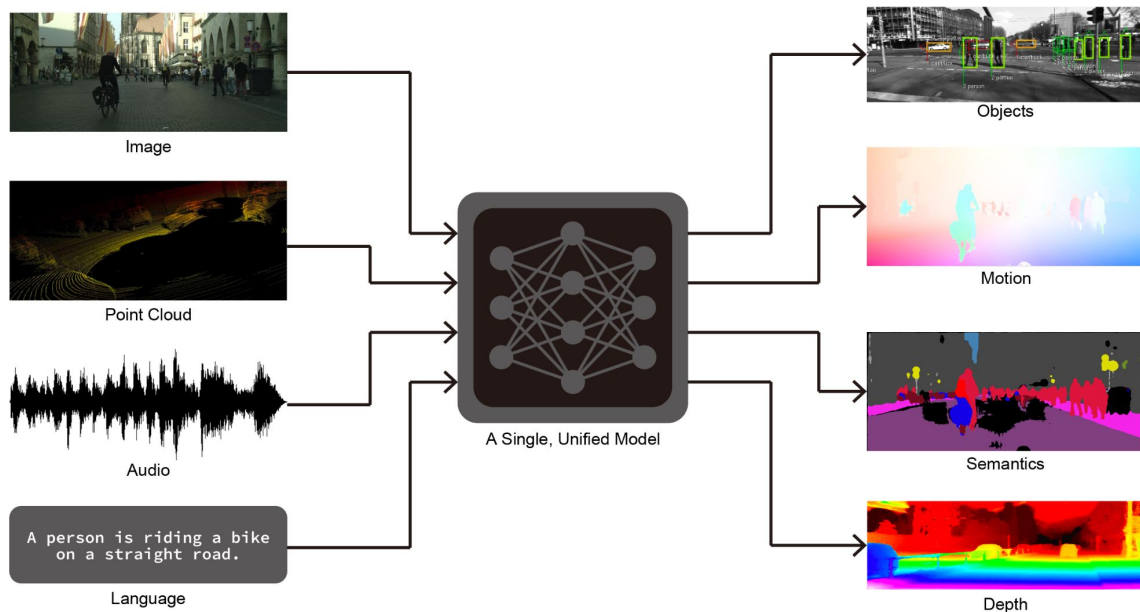
- **Inductive bias (formal definition):** a set of assumptions that the learner uses to predict outputs given input that has not observed
- Many neural network architectures rely on certain type of inductive biases
 - Mostly depends on locality and invariance of data
 - Example1: Convolution exploits 2D grid data structure and translation invariance
 - Example 2: RNN exploits 1D chain data structure and first-order Markov assumption
- These assumptions are generally useful to design efficient models

Inductive bias in neural networks

- **Inductive bias (formal definition):** a set of assumptions that the learner uses to predict outputs given input that has not observed
- Many neural network architectures rely on certain type of inductive biases
 - Mostly depends on locality and invariance of data
 - Example1: Convolution exploits 2D grid data structure and translation invariance
 - Example 2: RNN exploits 1D chain data structure and first-order Markov assumption
- These assumptions are generally useful to design efficient models
- The problem is that they often not generalize across different data modalities

Towards versatile neural networks

- We want to design a single model that handles many different data / tasks
- What is the most important requirement for this?




Content

- Revisiting attention and Transformers for versatile architectures
- Extending Transformers to heterogeneous data and tasks

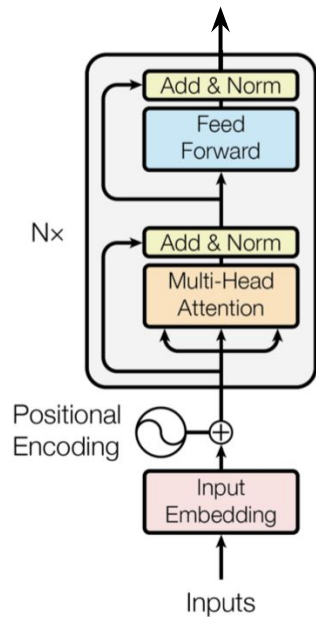
Revisit: Dot-product attention

$$\text{Attn}(\mathbf{x}W_q, \mathbf{x}W_k, \mathbf{x}W_v) = \text{softmax} \left(\frac{\mathbf{x}W_q(\mathbf{x}W_k)^T}{\sqrt{d}} \right) \mathbf{x}W_v$$



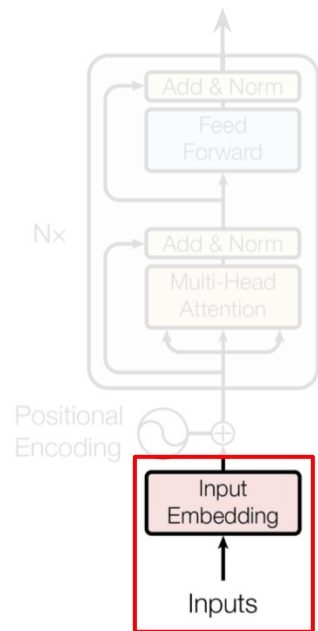
Revisit: Transformers (encoder only)

- Input is tokenized $\mathbf{x} \in \mathbb{R}^{n \times d}$



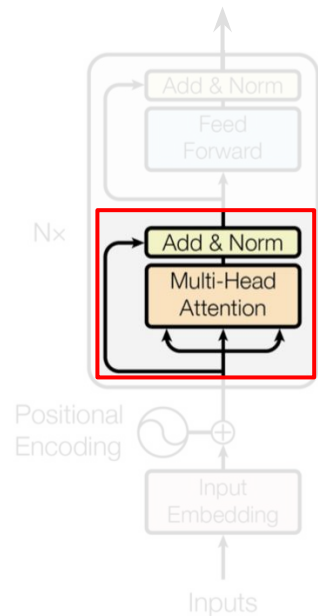
Revisit: Transformers (encoder only)

- Input is tokenized $\mathbf{x} \in \mathbb{R}^{n \times d}$
- Each token is encoded into continuous embedding
 $\mathbf{h}_i = \text{FF}(\mathbf{x}_i)$



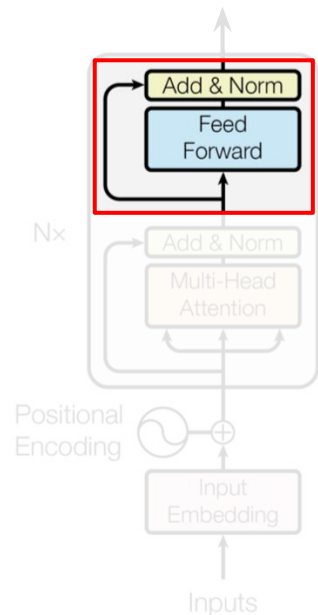
Revisit: Transformers (encoder only)

- Input is tokenized $\mathbf{x} \in \mathbb{R}^{n \times d}$
- Each token is encoded into continuous embedding
 $\mathbf{h}_i = \text{FF}(\mathbf{x}_i)$
- The tokens are then applied contexture encoding based on multi-head attention with skip connection
 $\mathbf{h}_i = \text{LayerNorm}(\mathbf{h}_i + \text{Attn}(\mathbf{h}_i, \mathbf{h}, \mathbf{h}))$



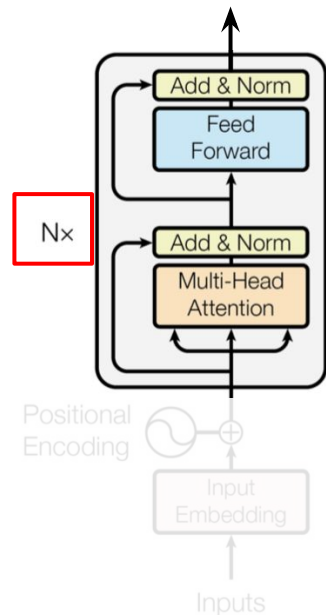
Revisit: Transformers (encoder only)

- Input is tokenized $\mathbf{x} \in \mathbb{R}^{n \times d}$
- Each token is encoded into continuous embedding
 $\mathbf{h}_i = \text{FF}(\mathbf{x}_i)$
- The tokens are then applied contexture encoding based on multi-head attention with skip connection
 $\mathbf{h}_i = \text{LayerNorm}(\mathbf{h}_i + \text{Attn}(\mathbf{h}_i, \mathbf{h}, \mathbf{h}))$
- It applies additional encode with skip connection
 $\mathbf{h}_i = \text{LayerNorm}(\mathbf{h}_i + \text{FF}(\mathbf{h}_i))$



Revisit: Transformers (encoder only)

- Input is tokenized $\mathbf{x} \in \mathbb{R}^{n \times d}$
- Each token is encoded into continuous embedding
 $\mathbf{h}_i = \text{FF}(\mathbf{x}_i)$
- The tokens are then applied contexture encoding based on multi-head attention with skip connection
 $\mathbf{h}_i = \text{LayerNorm}(\mathbf{h}_i + \text{Attn}(\mathbf{h}_i, \mathbf{h}, \mathbf{h}))$
- It applies additional encode with skip connection
 $\mathbf{h}_i = \text{LayerNorm}(\mathbf{h}_i + \text{FF}(\mathbf{h}_i))$
- The output is recursively served as the next input

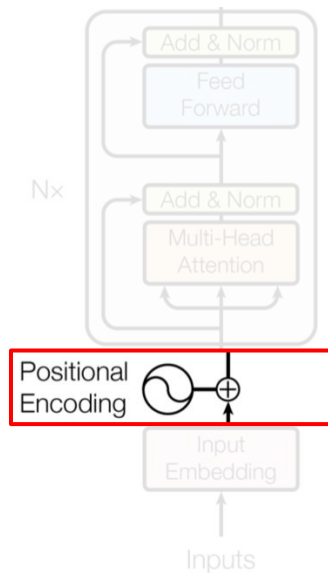


Revisit: Transformers (encoder only)

- When encoding tokens, we may add extra embedding that encodes **position** of the token.
- We can use any periodic function for positional embedding, or even learn how to encode position


$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$




Transformer: performance

- The original Transformer paper was developed for machine translation
- It turned out that Transformer scales surprisingly well with large-scale dataset
- Most innovations in large-scale models are based on Transformer architecture



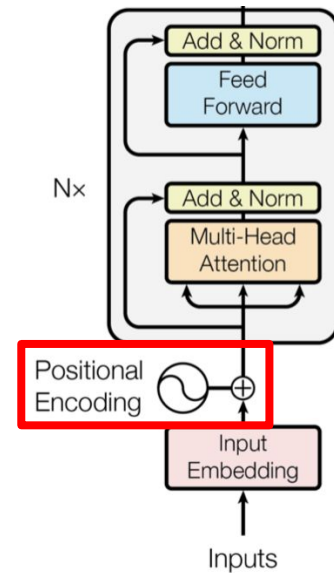
Transformer: analysis

- What are the inductive biases in the Transformers?
 - Data can be represented as a set of “tokens”



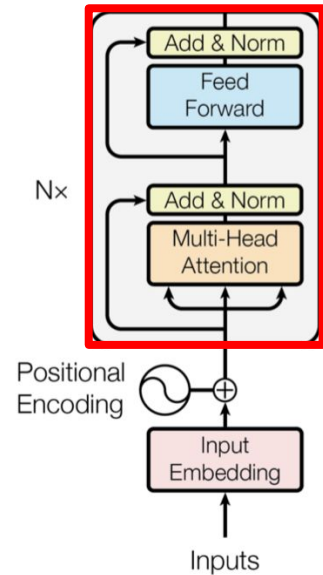
Transformer: analysis

- What are the inductive biases in the Transformers?
 - Data can be represented as a set of “tokens”
 - Position of the tokens can be represented by the “positional encoding”




Transformer: analysis

- What are the inductive biases in the Transformers?
 - Data can be represented as a set of “tokens”
 - Position of the tokens can be represented by the “positional encoding”
 - Patterns in the data can be discovered by the “pairwise relational reasoning”




Transformer: analysis


- What are the inductive biases in the Transformers?
 - Data can be represented as a set of “tokens”
 - Position of the tokens can be represented by the “positional encoding”
 - Patterns in the data can be discovered by the “pairwise relational reasoning”
- Are these inductive biases generally applicable across data?




Language



Image

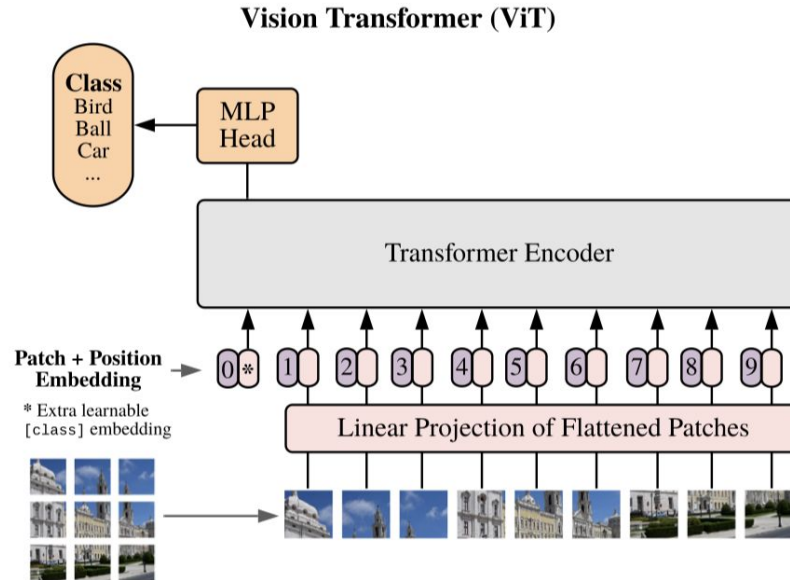


Video



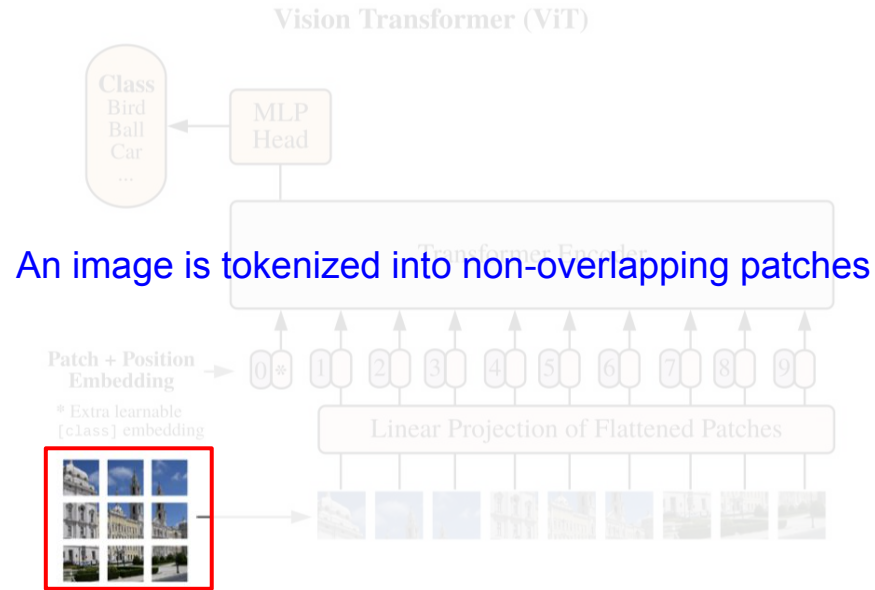
Transformers for vision

- Vision Transformer for image classification



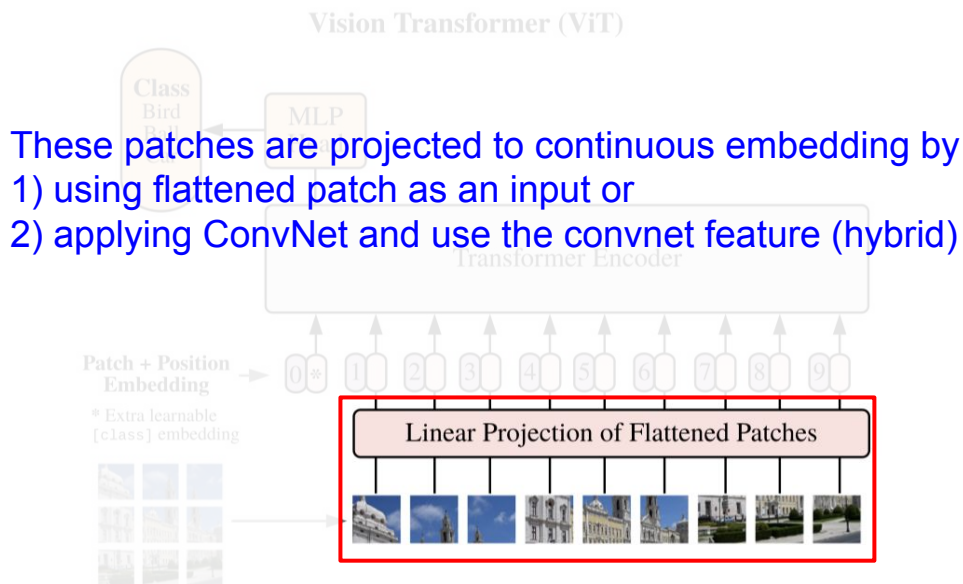
Transformers for vision

- Vision Transformer for image classification



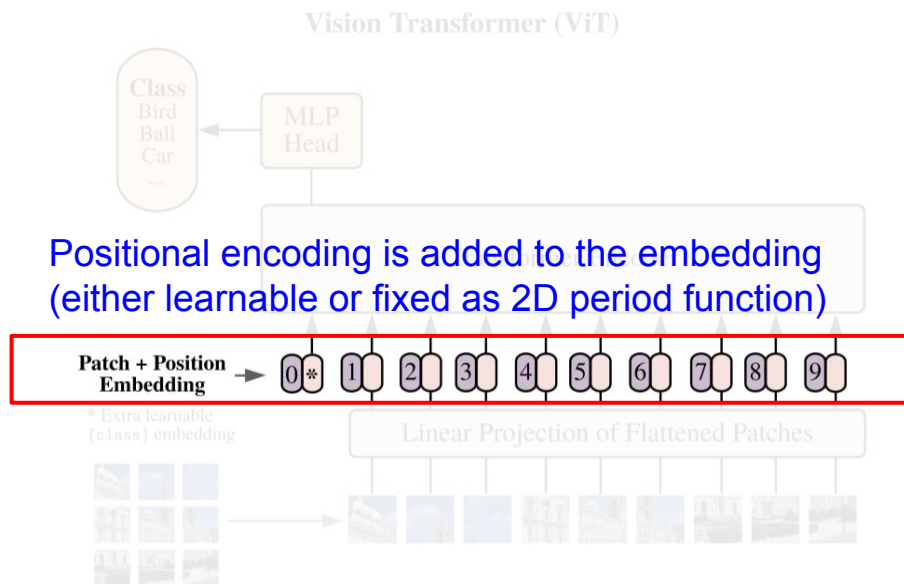
Transformers for vision

- Vision Transformer for image classification



Transformers for vision

- Vision Transformer for image classification



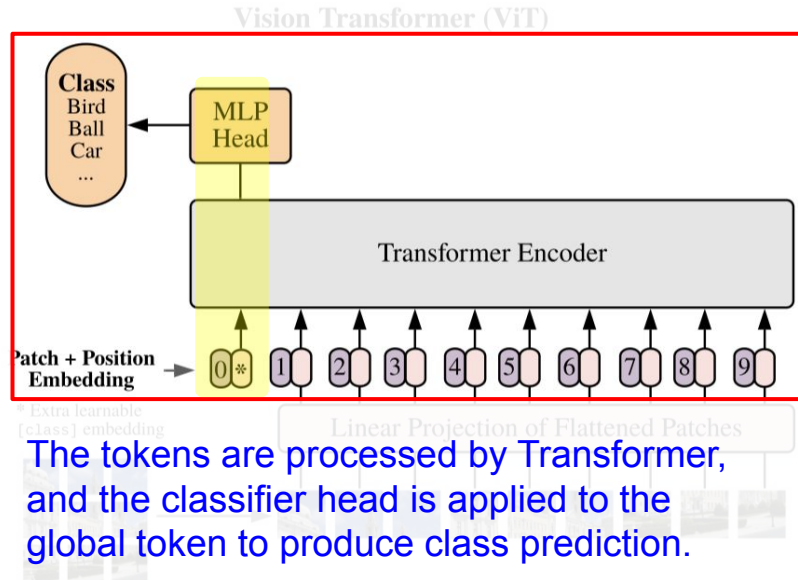
Transformers for vision

- Vision Transformer for image classification




Transformers for vision

- Vision Transformer for image classification



Transformers for vision

- Vision Transformer for image classification



Transformers for vision


- Vision Transformer outperforms the ResNet with large-scale pre-training

	Ours-JFT (ViT-H/14)	Ours-JFT (ViT-L/16)	Ours-I21k (ViT-L/16)	BiT-L (ResNet152x4)	Noisy Student (EfficientNet-L2)
ImageNet	88.55 ± 0.04	87.76 ± 0.03	85.30 ± 0.02	87.54 ± 0.02	88.4/88.5*
ImageNet ReaL	90.72 ± 0.05	90.54 ± 0.03	88.62 ± 0.05	90.54	90.55
CIFAR-10	99.50 ± 0.06	99.42 ± 0.03	99.15 ± 0.03	99.37 ± 0.06	—
CIFAR-100	94.55 ± 0.04	93.90 ± 0.05	93.25 ± 0.05	93.51 ± 0.08	—
Oxford-IIIT Pets	97.56 ± 0.03	97.32 ± 0.11	94.67 ± 0.15	96.62 ± 0.23	—
Oxford Flowers-102	99.68 ± 0.02	99.74 ± 0.00	99.61 ± 0.02	99.63 ± 0.03	—
VTAB (19 tasks)	77.63 ± 0.23	76.28 ± 0.46	72.72 ± 0.21	76.29 ± 1.70	—
TPUv3-core-days	2.5k	0.68k	0.23k	9.9k	12.3k

Transformers for vision


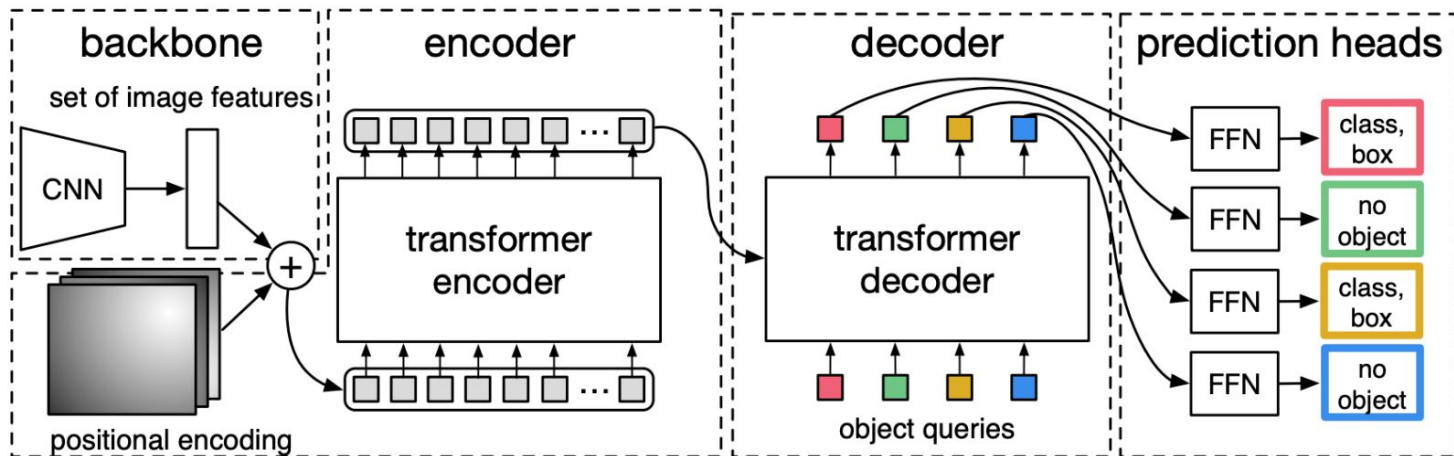
- ViT scales well with the larger data

The benefit from hybrid architecture vanishes as it is learned from larger data



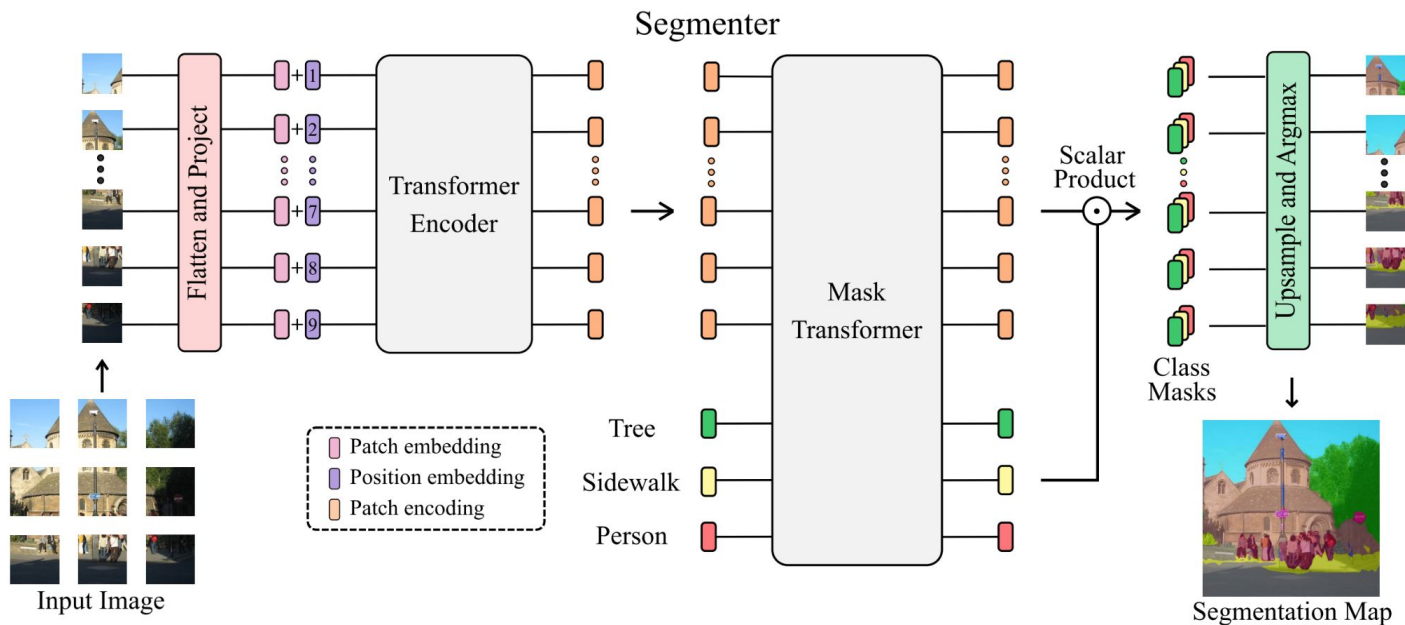
Transformers for structured vision prediction

- Transformer + object detection



Transformers for structured vision prediction

- Transformer + semantic segmentation




Can we treat images as words?

- Vision Transformers (ViT) treat patches as tokens
- However, these patches are continuous while words are discrete
- If we can discretize the images (patches) as words, then there is no difference in language and vision


Discretizing image representation

- Basic idea
 - Learn an autoencoder (or VAE)





Discretizing image representation

- Basic idea
 - Learn an autoencoder (or VAE)
 - Collect all learned (patch) embeddings of training images





Discretizing image representation

- Basic idea
 - Learn an autoencoder (or VAE)
 - Collect all features of training images
 - Then apply the clustering in the embedding space, which will give us cluster centers



Discretizing image representation


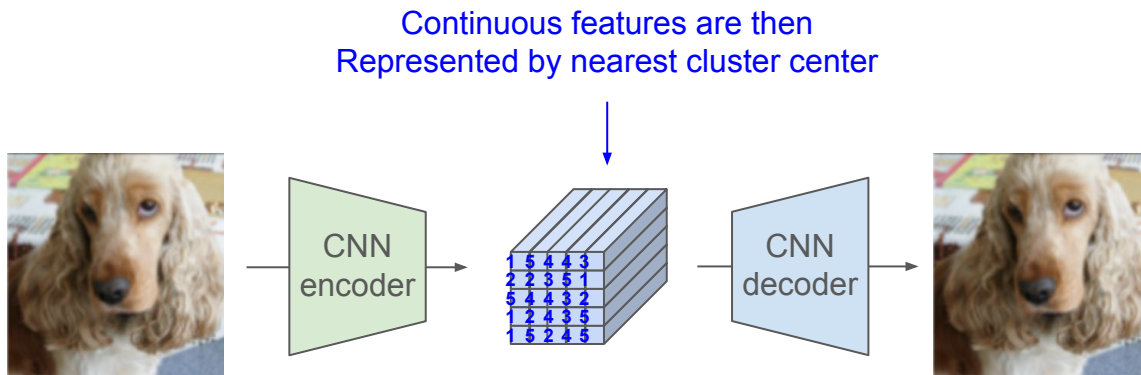
- Basic idea
 - Learn an autoencoder (or VAE)
 - Collect all features of training images
 - Then apply the clustering in the embedding space, which will give us cluster centers
 - Assigning the unique indices to the cluster centers, every continuous embeddings can be assigned with discrete index by associating them to nearest cluster centers



Discretizing image representation

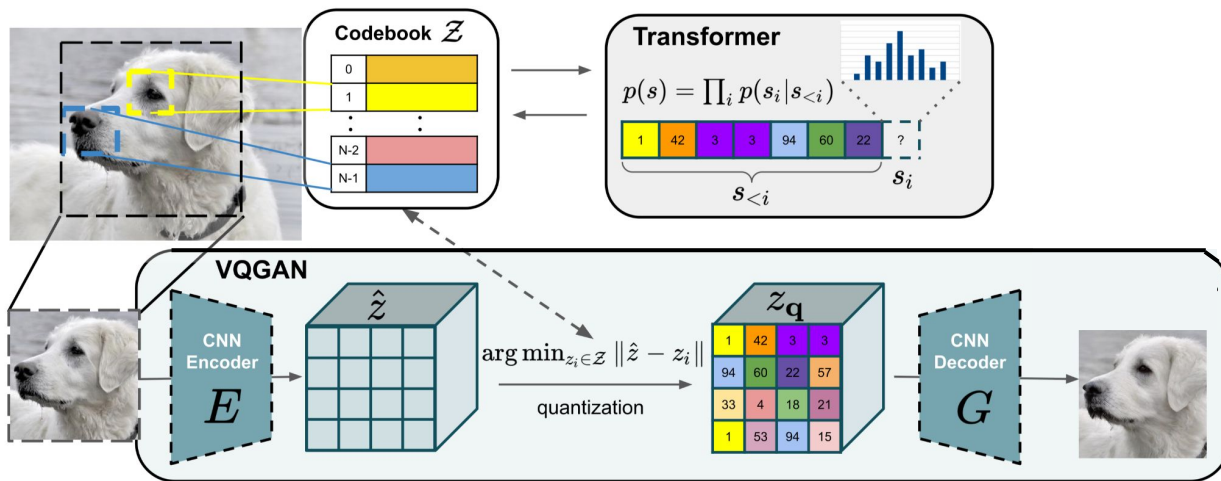
- Basic idea

- Learn an autoencoder (or VAE)
- Collect all features of training images
- Then apply the clustering in the embedding space, which will give us cluster centers
- Assigning the unique indices to the cluster centers, every continuous embeddings can be assigned with discrete index by associating them to nearest cluster centers



Vector quantization

- Learn to cluster (quantize) the features end-to-end with autoencoding

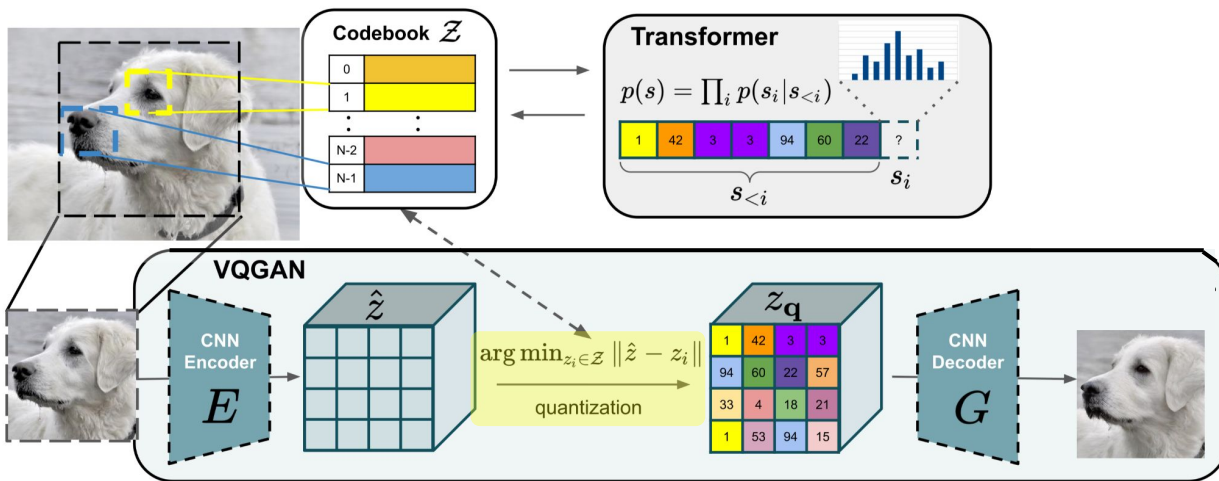


Vector quantization

- Learn to cluster (quantize) the features end-to-end with autoencoding

$$z_q = \mathbf{q}(\hat{z}) := \left(\arg \min_{z_k \in \mathcal{Z}} \|\hat{z}_{ij} - z_k\| \right) \in \mathbb{R}^{h \times w \times n_z}$$

Find the nearest codebook (cluster center) and replace the feature with it

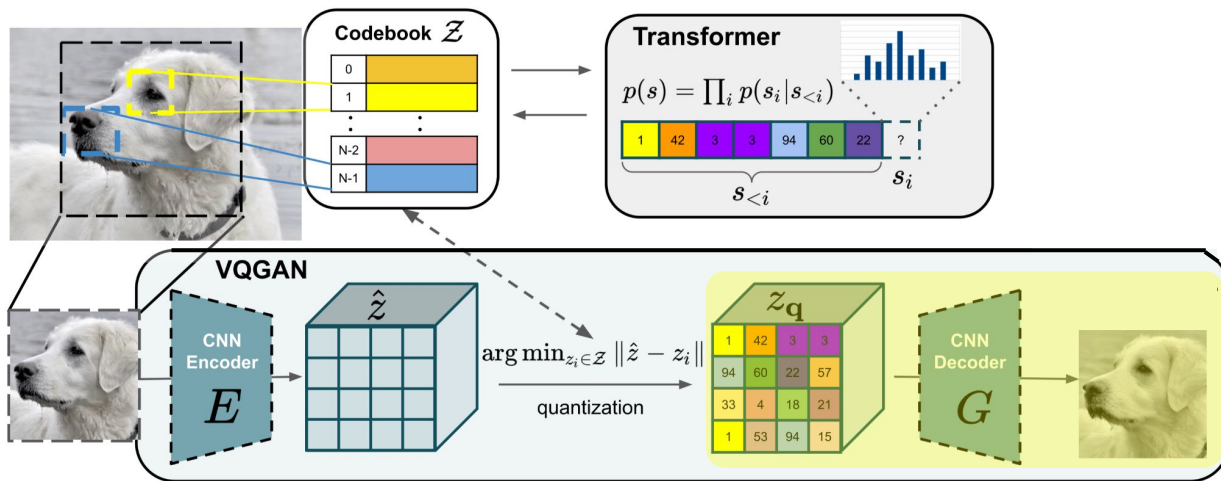


Vector quantization

- Learn to cluster (quantize) the features end-to-end with autoencoding

$$\hat{x} = G(z_q) = G(q(E(x)))$$

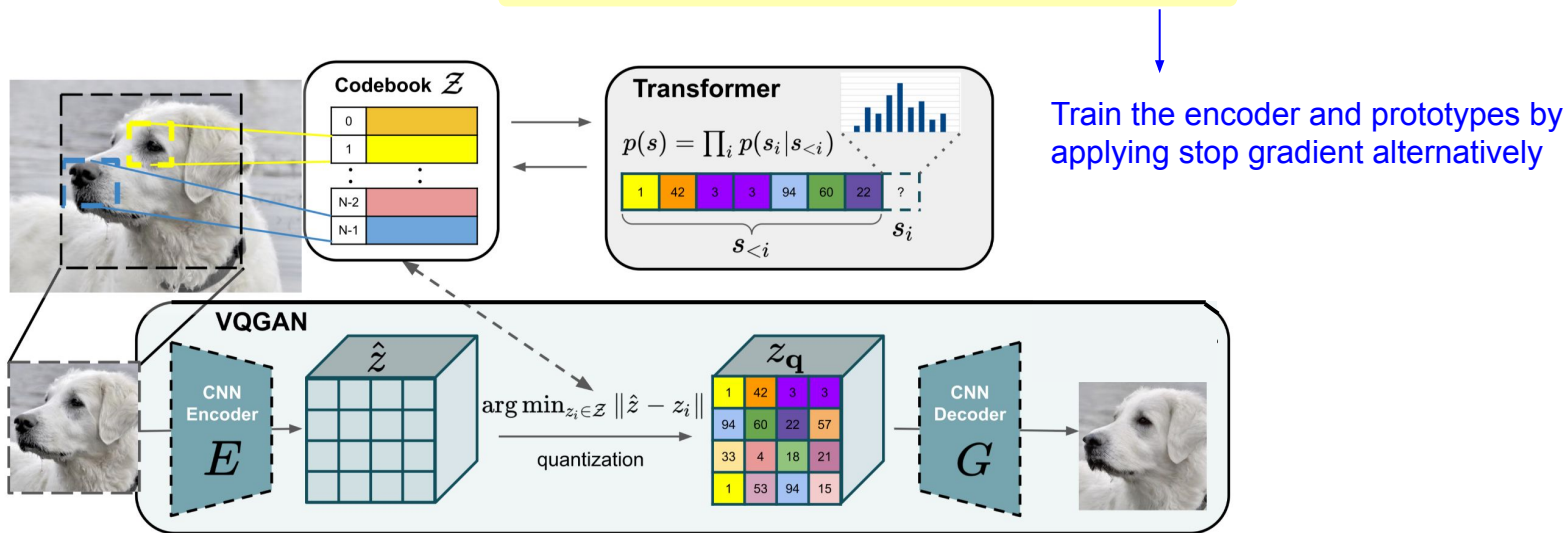
The output image is obtained by decoding the quantized features



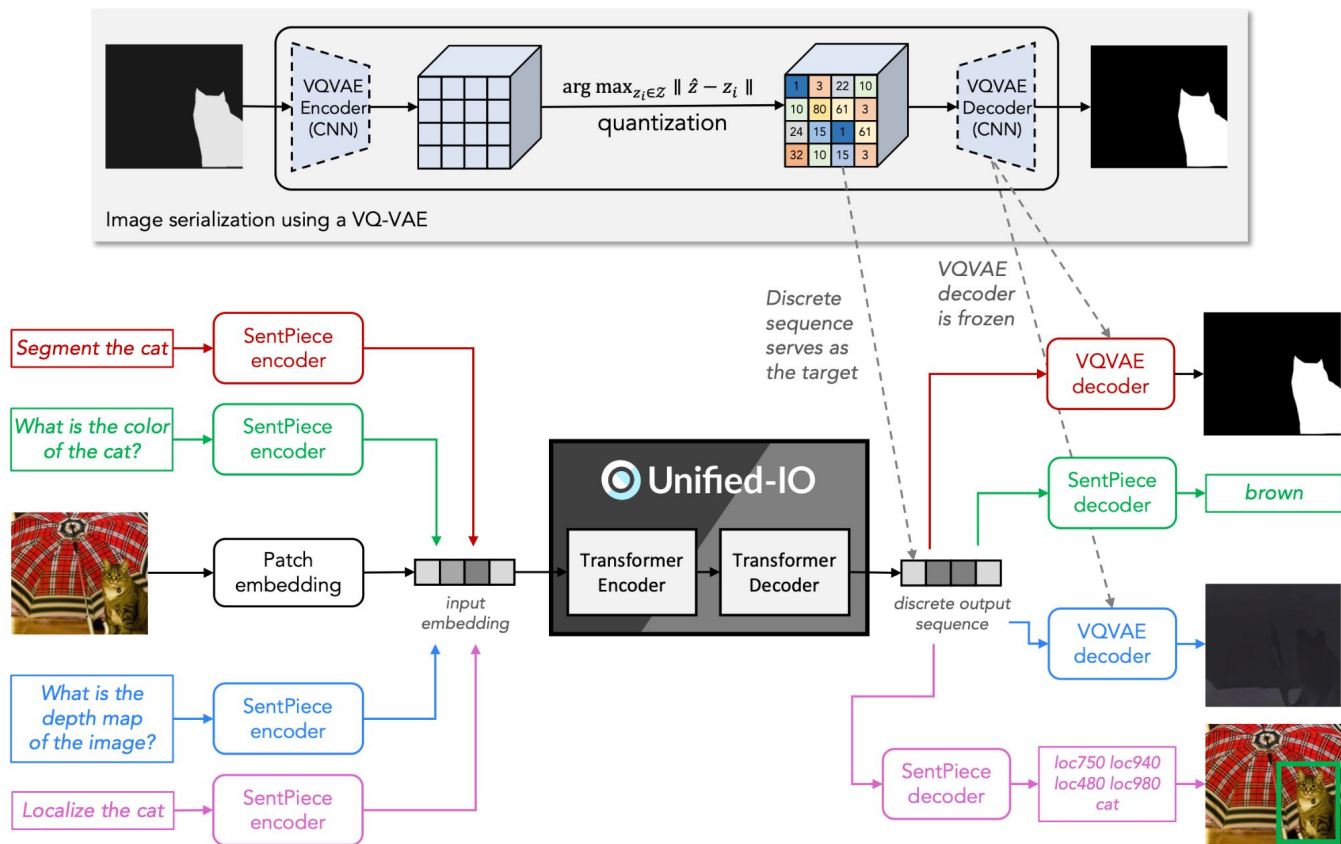
Vector quantization

- Learn to cluster (quantize) the features end-to-end with autoencoding

$$\mathcal{L}_{\text{VQ}}(E, G, \mathcal{Z}) = \|x - \hat{x}\|^2 + \|\text{sg}[E(x)] - z_{\mathbf{q}}\|_2^2 + \beta \|\text{sg}[z_{\mathbf{q}}] - E(x)\|_2^2$$



Transformers for universal vision learner



Transformer for language, vision, and RL agents

